

Quick Announcements

Update on Case Study

- 1 student per team sends, via canvas message to the TAs, the following information:
 - List of team members
 - Topic chosen by the team
 - Description (1-2 sentences) of each member's part

“warm up” PAs from previous slide PDF

- All are trivial except the “It's meta” one
- Post on piazza if you have questions about any of them



It's meta.

- We want to assign to a variable called `META` the value `$RANDOM`. Please note, we don't want to assign a random number but the value `$RANDOM` as a string. **How do you do that?**
- We then want to display a message that states: Your lucky number is 112233. The value 112233 will be a random number, different every time we re-execute the line, and generated by using the variable `META`

Why doesn't the following work?

How do we make it work?

- Hint: check out the doc for `eval`

```
$ echo $( echo $META)
$RANDOM

$ echo $(echo $( echo $META))
$RANDOM
```

M03

The UNIX Way

Menu for this module

T1	Redirections & Piping	It is possible, and useful, to direct the text output of a given command line tool so that it is accepted as input of another. There "redirections" are part of what makes command line useful to automate tasks.
T2	Filters	For the above to be really leveraged, we need tools which are designed to work line per line on the data fed to them from the keyboard, a file, or another tool's output. Enter the filters....

M3T1.1


Bash Redirections Rudiments



<https://youtu.be/yV0wCSm2YBY>

Bash Redirections - Rudiments

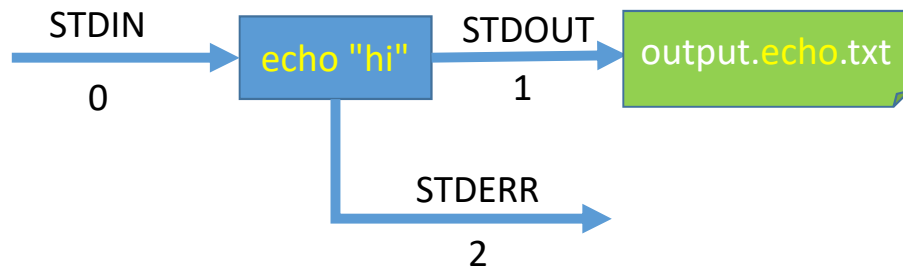
- Pipelines of information between processes
- Each process has 3 file descriptors



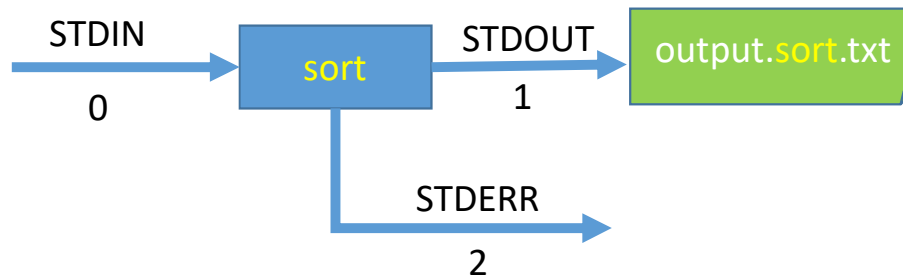
```
$ sort
one
two
three
^D
one
three
two
```

```
$ sort -l
sort: invalid option -- 'l'
Try 'sort --help' for more
information.
```

Redirecting **STDOUT** to file




```
$ echo "hi" > output.echo.txt
```



```
$ sort 1> output.sort.txt
One
Two
Three
^D
```

NOTE – STDERR is still displaying to the console with a `sort -l`


Redirecting **STDERR** to file



```
$ sort -l > logs.txt  
$ cat logs.txt  
  
$ sort -l 1> logs.txt  
$ cat logs.txt  
  
$ sort -l 2> logs.txt  
$ cat logs.txt
```

- `>` and `1>` do not work for this
- We need a new syntax: **`2>`**

Redirecting **both** STDOUT & STDERR to **2** files



```
$ sort -l 2> sort.log 1> sort.out
$ cat sort.log
$ cat sort.out
```


```
$ # does order matter?
```

```
$ echo "hi" 2> echo.log 1> echo.out
$ cat echo.log
$ cat echo.out
```

```
$ echo "hi" 1> echo.out 2> echo.log
$ cat echo.log
$ cat echo.out
```

```
$ # order does not matter
```

What about **STDIN**?



```
$ cat > somedata.txt
We are
Making a new file
With cat redirecting its output
^D

$ sort < somedata.txt
$ sort 0< somedata.txt

$ sort -l < somedata.txt
sort: invalid option -- 'l'
Try 'sort --help' for more information.
```

Let's “Fire Everything!!!” that we got so far



```
$ sort < somedata.txt 2> sort.log 1> sort.out  
$ cat sort.log  
$ cat sort.out
```

/dev/null

```
$ cat /dev/null
```

```
$ sort -l 2> /dev/null
```

```
$ echo "into the black hole!" 1> /dev/null
```

```
$ echo "into the black hole!" > /dev/null
```

M3T1.2

Bash Redirections - Appending



<https://youtu.be/NknAcynU4sM>

Accumulating **STDERR** msgs to log file

> 1> 2> means overwrite
>> 1>> 2>> means **append**

```
sort -l 2> sort.err  
sort -l 2> sort.err  
cat sort.err
```

```
# only 2nd message is there
```

```
sort -l 2>> sort.err  
sort -l 2>> sort.err  
cat sort.err
```

```
# both messages are there
```

Trying now with **STDOUT**

- Works the same way

```
echo "hello" > echo.out
cat echo.out

echo "hello" 1> echo.out
echo "world" 1> echo.out
cat echo.out

# only 2nd message is there

rm echo.out

echo "hello" >> echo.out
echo "world" >> echo.out
cat echo.out

# both messages are there
```

Bash **option** to control that behavior

- If we redirect with overwrite to an existing file, I want to **prevent** losing all existing data

```
set -o noclobber

echo "hello world" > echo.out
Bash: echo.out: cannot overwrite
existing file

cat echo.out #all is still there

set +o noclobber

echo "hello world" > echo.out

cat echo.out #data is gone!!!!
```


Overriding the safety option

- If we redirect with overwrite to an existing file, I want to **prevent** losing all existing data
- What if I want an **exception** for just one command?
- Works with STDOUT, try it with STDERR it works as well

```
set -o noclobber

echo "hello world" > echo.out
Bash: echo.out: cannot overwrite existing file
cat echo.out #all is still there

echo "replacing data by this" >| echo.out
# override noclobber option
cat echo.out
```


M3T1.3 Bash Redirections

~~Crossing~~ Merging the Streams



<https://youtu.be/92-YIUXGh68>


Sometimes we want STDOUT & STDERR to be redirected to the **same location**



What happened?
Why didn't it work?

```
sort 1> sort.log 2> sort.log
Bash: sort.log: cannot overwrite
existing file
```

Here is the **explanation...**



? How do we fix it?


```
sort 1> sort.log 2> sort.log
```

```
Bash: sort.log: cannot overwrite  
existing file
```

```
# 1st redirection created log file
```

```
# 2nd redirection targeted the same file
```

This is how we fix it...



It works when we use STDOUT
or STDERR... what about **AND**?

```
sort 1> sort.log 2> sort.log
```

```
Bash: sort.log: cannot overwrite  
existing file
```

```
# 1st redir created log file
```

```
# 2nd redir targeted the same file
```

```
# if we get rid of noclobber..
```

```
set +o noclobber
```

```
sort 1> sort.log 2> sort.log
```

```
cat sort.log
```

```
# YAY! the output is in the log
```

```
sort -l 1> sort.log 2> sort.log
```

```
cat sort.log
```

```
# YAY! error message is in the log
```

What if our command generates both STDOUT and STDERR?

- Testing with `myprog.sh`

```
#!/bin/bash
```

```
>&2 echo "This is for STDERR"  
echo "This is for STDOUT"
```

```
./myprog.sh
```

```
This is something for STDOUT
```

```
This is something for STDERR
```

```
./myprog.sh 2>/dev/null
```

```
This is something for STDOUT
```

```
./myprog.sh 1>/dev/null
```

```
This is something for STDERR
```

```
./myprog.sh 1>/dev/null 2>/dev/null
```


```
# → able to redirect BOTH to /dev/null
```

```
# But
```

```
# what if we redirect to something else?
```

Let's **redirect to a file!**

Problem: 1 of the messages is lost!



```
./myprog.sh 1> log 2> log
cat log
This is something for STDERR

# only 1 of the messages made it to log
```

? Why did we lose 1 of the 2 messages?

- We turned off noclobber
- So now we have clobbering 😊

Solution: new syntax

- New Syntax!!!

2>&1

“redirecting FD#2 to where FD#1
is pointing at right now”

```
./myprog.sh 2>&1  
This is something for STDOUT  
This is something for STDERR
```


Solution: new syntax

- New Syntax!!!

2>&1

“redirecting FD#2 to where FD#1
is pointing at right now”

```
./myprog.sh 2>&1
```

```
This is something for STDOUT
```

```
This is something for STDERR
```

```
# ok but how do I check these were both  
# sent to STDOUT?
```

Solution: new syntax

- New Syntax!!!

2>&1

“redirecting FD#2 to where FD#1 is pointing at right now”

- New bug!!!! ;p

```
./myprog.sh 2>&1
```

```
This is something for STDOUT
```

```
This is something for STDERR
```

```
# ok but how do I check these were both  
# sent to STDOUT?
```

```
# Redirecting 2 to null should have no  
# effects since it is already on 1
```

Solution: new syntax

- New Syntax!!!

2>&1

“redirecting FD#2 to where FD#1 is pointing at right now”

- New bug!!!! ;p

```
./myprog.sh 2>&1
```

```
This is something for STDOUT
```

```
This is something for STDERR
```

```
# ok but how do I check these were both  
# sent to STDOUT?
```

```
# Redirecting 2 to null should have no  
# effects since it is already on 1
```

```
./myprog.sh 2>&1 2>/dev/null
```

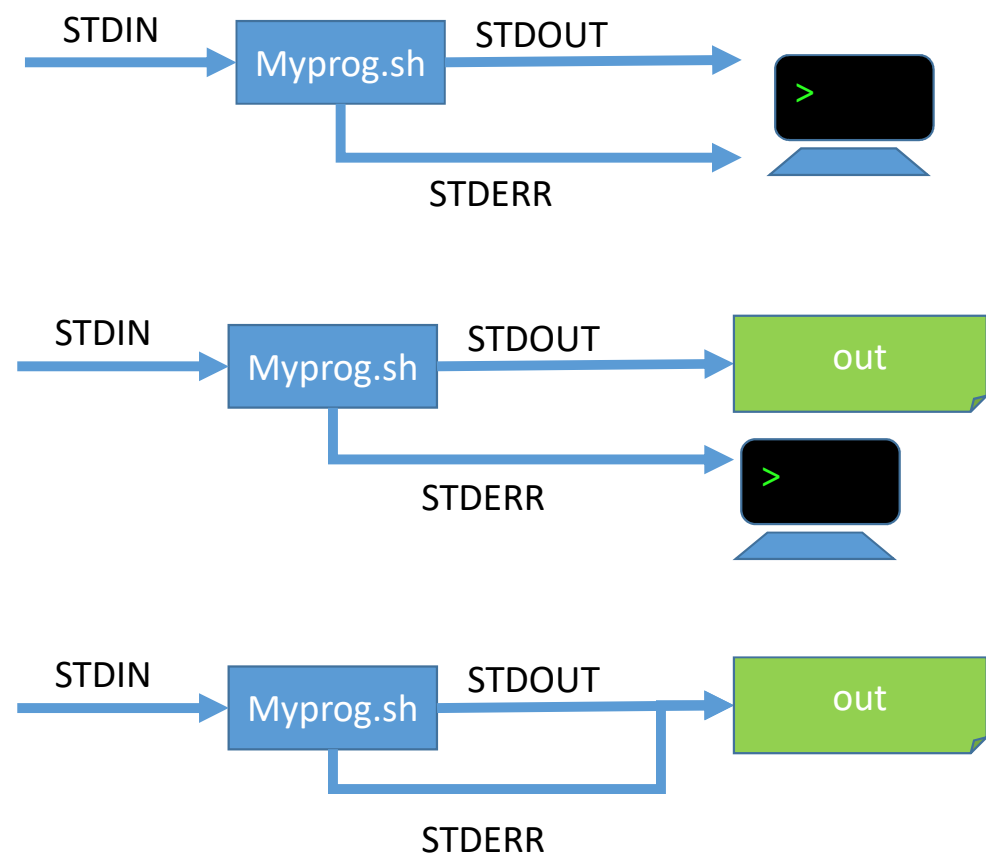
```
This is something for STDOUT
```

```
# WAIT!? WHAT????
```

```
# FD#2 was set to FD#1 destination
```

```
# but then we reset FD#2 to /dev/null
```

Interlude (1/2) – What happens if we re-redirect?




```
./myprog.sh
This is something for STDOUT
This is something for STDERR
```

```
./myprog.sh > out
This is something for STDERR
cat out
This is something for STDOUT
```

```
./myprog.sh > out 2>&1
cat out
This is something for STDOUT
This is something for STDERR
```

Interlude (2/2) – What if we reverse the order?




```
./myprog.sh  
This is something for STDOUT  
This is something for STDERR
```

```
./myprog.sh 2>&1  
This is something for STDOUT  
This is something for STDERR
```

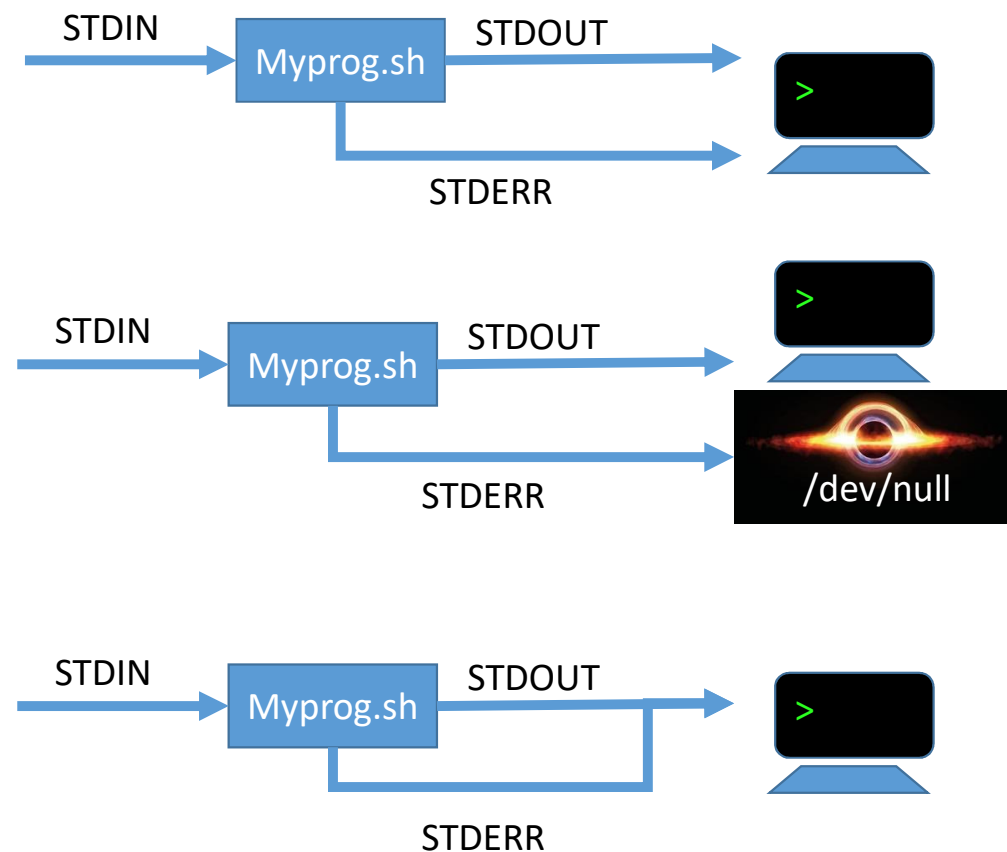
```
./myprog.sh 2>&1 > out  
This is something for STDERR  
cat out  
This is something for STDOUT
```

Back to our redirection to /dev/null



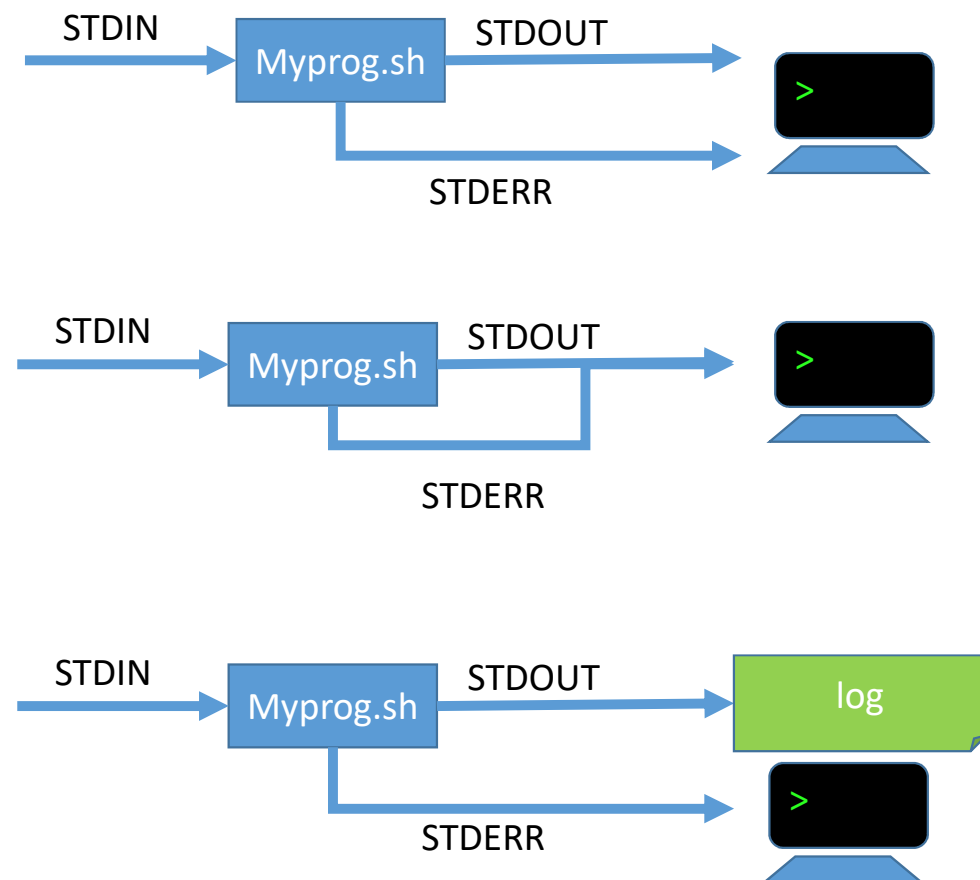
```
./myprog.sh 2>&1 2>/dev/null  
This is something for STDOUT
```

Let's switch it around 1 more time!
(To make sure that we understand)






```
./myprog.sh 2>/dev/null 2>&1  
This is something for STDOUT  
This is something for STDERR
```

What if we `2>&1` but THEN change `1>...`



```
./myprog.sh 2>&1 > log  
This is something for STDERR  
  
# only 1 msg in the file!  
# STDERR still on console!
```


What if we `2>&1` but THEN change `1>...`

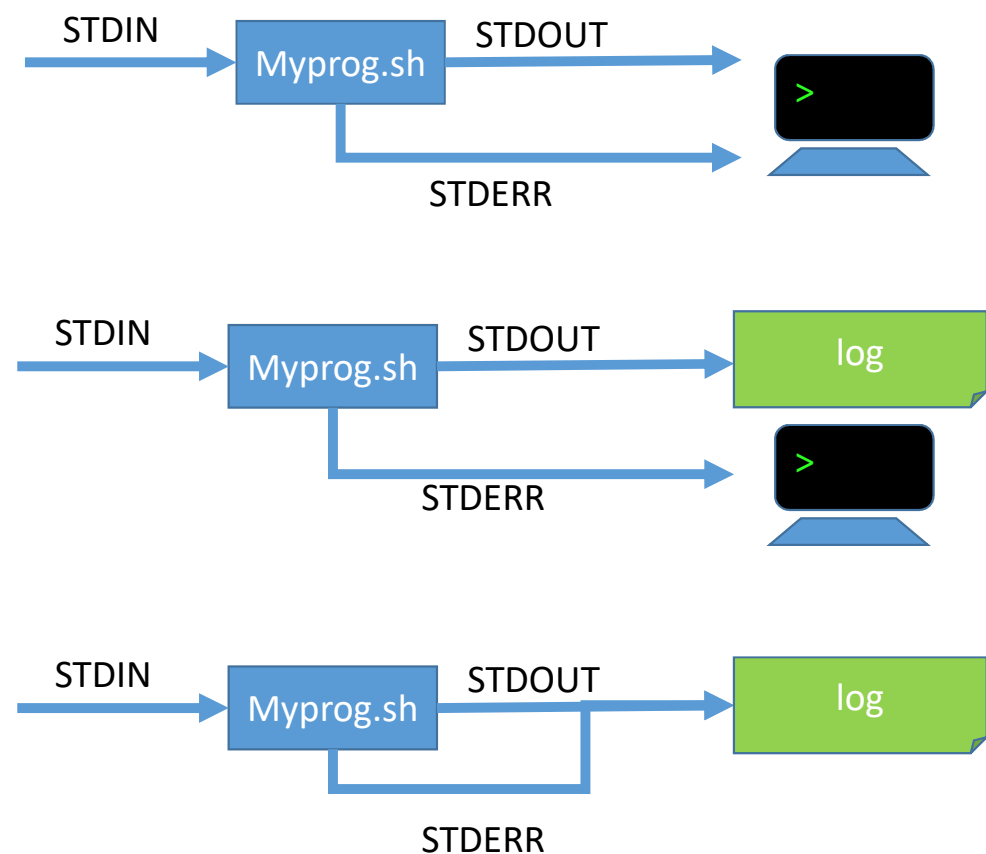


```
./myprog.sh 2>&1 > log  
This is something for STDERR
```

```
# only 1 msg in the file!  
# STDERR still on console!
```

← How do we get this instead?

What if we `2>&1` but THEN change `1>...`



```
./myprog.sh 2>&1 > log  
This is something for STDERR
```

```
# only 1 msg in the file!  
# STDERR still on console!
```

← How do we get this instead?

```
./myprog.sh > log 2>&1
```


M3T1.4

Bash Redirections – Swapping STDOUT & STDERR



<https://youtu.be/0kY9nFO1078>


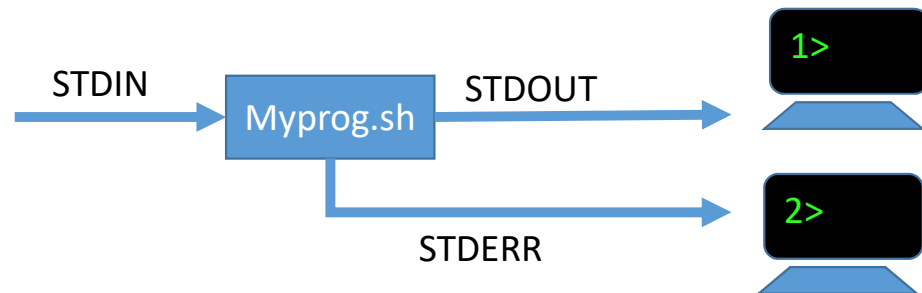
Challenge: how to swap STDOUT / STDERR



Let's start by only using what we learned so far



- > 1> 2>
- >> 1>> 2>>
- 1>&2 2>&1

Trying `2>&1`



We **lost STDERR** target,
whether it was the
console or a file or a
pipe...

Trying 1>&2



We **lost** **STDOUT** target, whether it was the console or a file or a pipe...

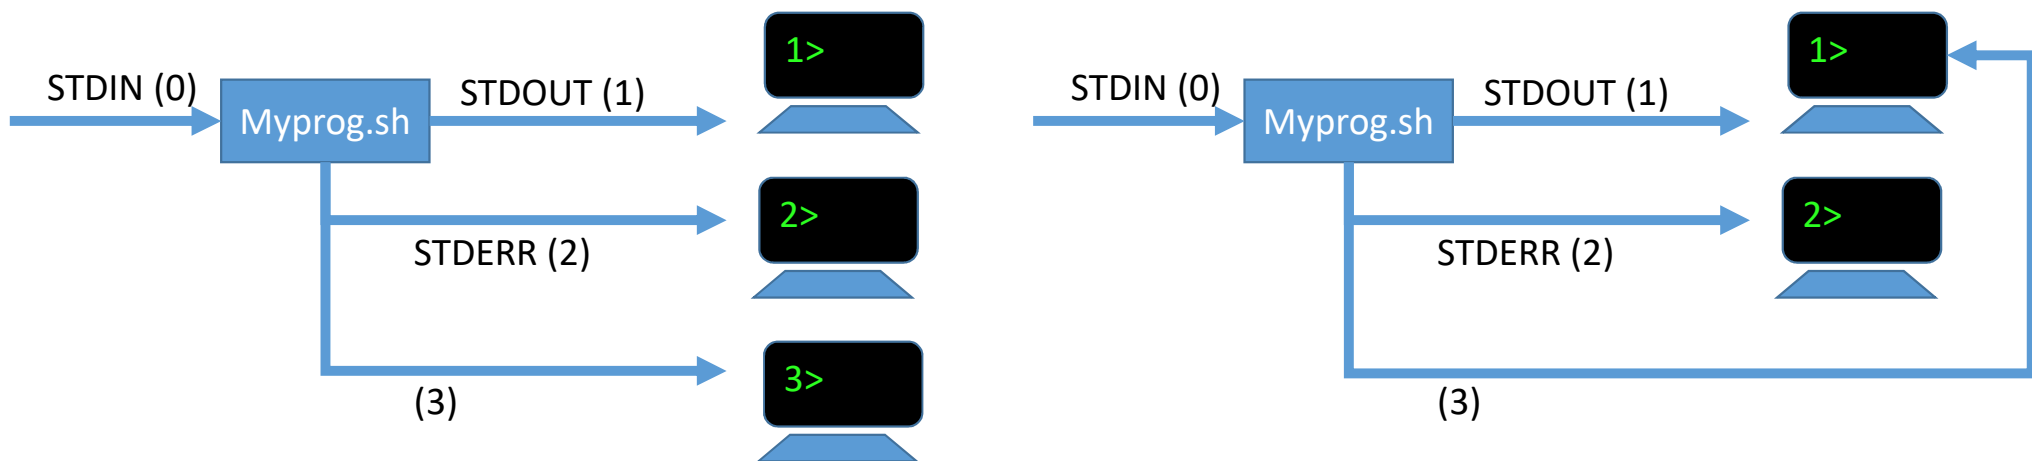
Problem is reminiscent of...

- Given two variables x and y , swap their contents.
- $x = y ; y = x$ \rightarrow we lost the value that was in x
- $y = x ; x = y$ \rightarrow we lost the value that was in y
- $tmp = y ; y = x ; x = tmp$ \rightarrow this works



Solution (1/3)

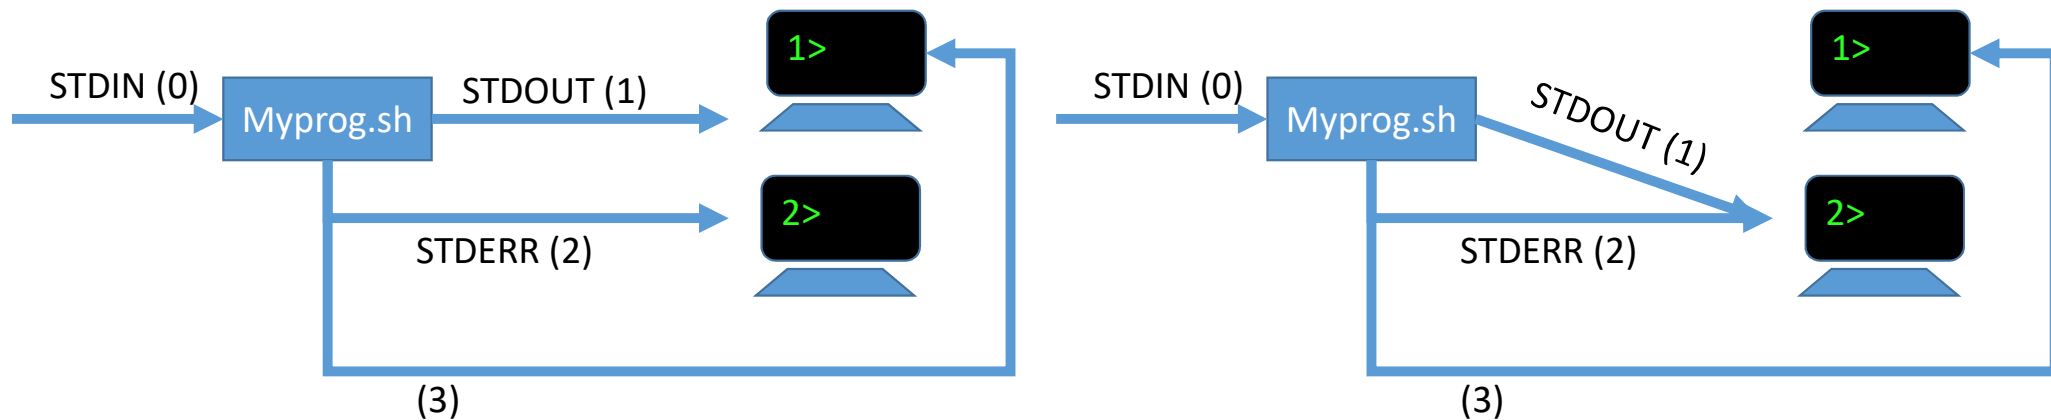
- 0, 1, and 2 are indexes in the file descriptors table
- There is a file descriptor at index 3 that we could use as TMP



- `3>&1` → saves FD #1 in FD #3

Solution (2/3)

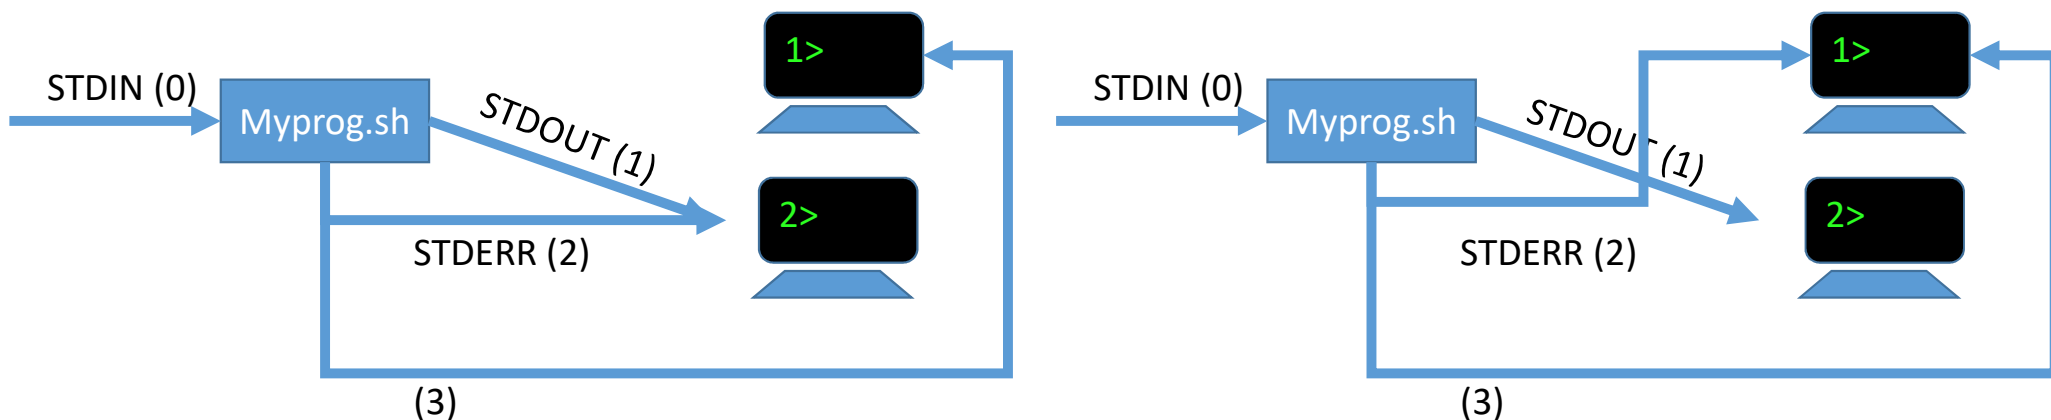
- 0, 1, and 2 are indexes in the file descriptors table
- There is a file descriptor at index 3 that we could use as TMP



- `3>&1` `1>&2`

Solution. (3/3)

- 0, 1, and 2 are indexes in the file descriptors table
- There is a file descriptor at index 3 that we could use as TMP



- `3>&1` `1>&2` `2>&3`

Let's apply this to myprog.sh

- Hardest thing we can do with what we learned so far so good to wrap up the topic
- Practical application: swapping two FDs
- **Generalizes** what we learned so far: 0,1,2 are not the only FDs available!



How do we verify that the swap really happened?

```
./myprog.sh 3>&1 1>&2 2>&3  
This is something for STDOUT  
This is something for STDERR  
  
# Cannot really tell that it worked  
# so we try something more...
```

Let's apply this to myprog.sh

- Hardest thing we can do with what we learned so far so good to wrap up the topic
- Practical application: swapping two FDs
- Generalizes what we learned so far: 0,1,2 are not the only FDs available!

```
./myprog.sh 3>&1 1>&2 2>&3
This is something for STDOUT
This is something for STDERR


# Cannot really tell that it worked
# so we try something more...

./myprog.sh > out 2> err 3>&1 1>&2 2>&3
cat out
This is something for STDERR
cat err
This is something for STDOUT
```

M3T1.5

Bash Redirections


– Piping



<https://youtu.be/pPEX5q6odFI>

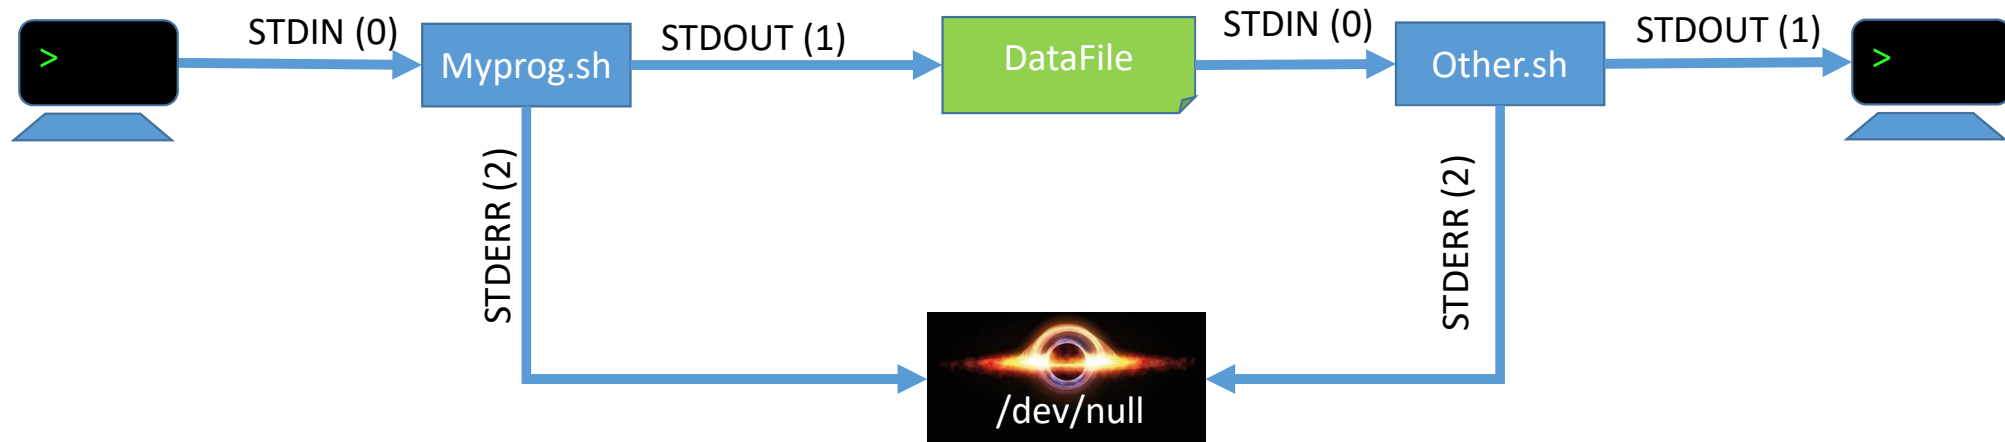
The story so far...

- Connecting FILES to STDOUT or STDERR
- What if we have two processes...



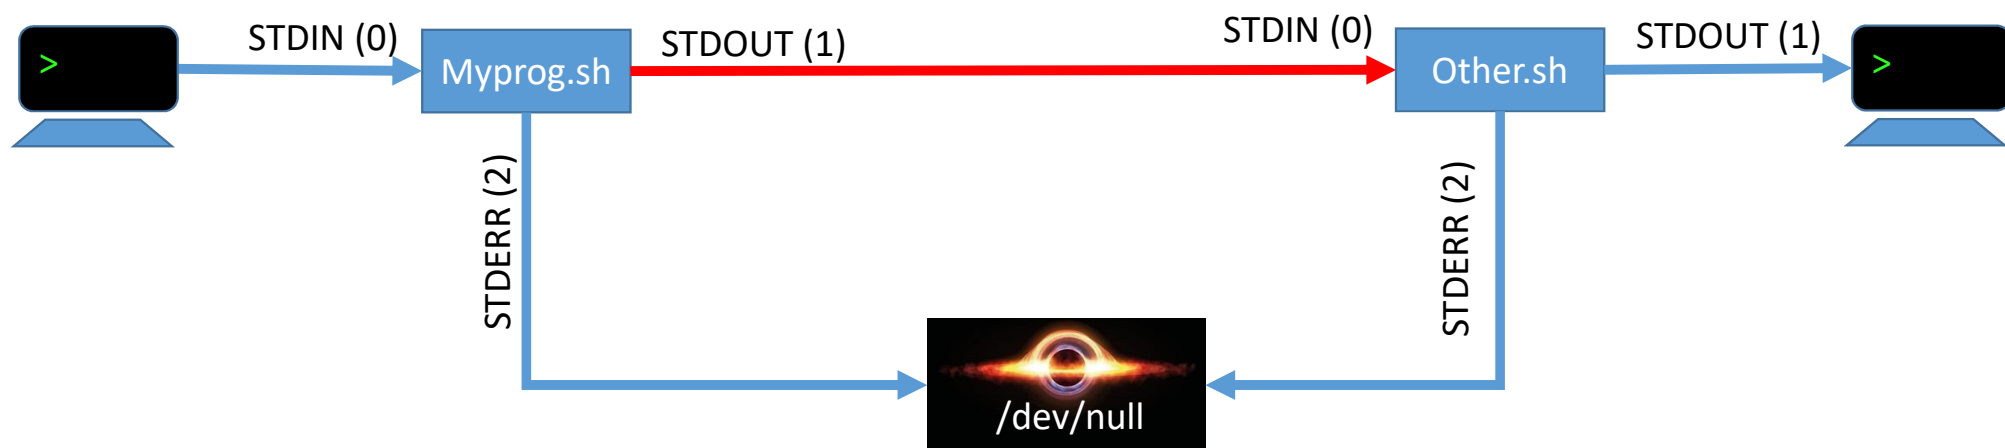
The story so far...

- Connecting FILES to STDOUT or STDERR
- What if we have two processes...
- Let's use **DataFile** to send data from 1st process to 2nd process



What about we skip using DataFile?


- Using **|** instead
- How do we illustrate this with simple programs?



Let's find some commands to illustrate this


```
cat
One
One
Two
Two
^D

cat > somedata.txt
One
Two
Three
^D
cat somedata.txt
One
Two
Three
```




We need a command that reads from STDIN (**kbd**) and displays on STDOUT (screen)

We also need a command that reads its input from STDIN and displays its output to STDOUT



```
sort
One
Two
Four
^D
Four
One
Two
```

cat & sort example:



```
cat | sort
```

```
One
```

```
Two
```

```
Four
```

```
^D
```

```
Four
```

```
One
```

```
Two
```

Introducing Tee. Earl Grey. Hot


```
cat | tee log | sort
```

```
One  
Two  
Four  
^D  
Four  
One  
Two
```

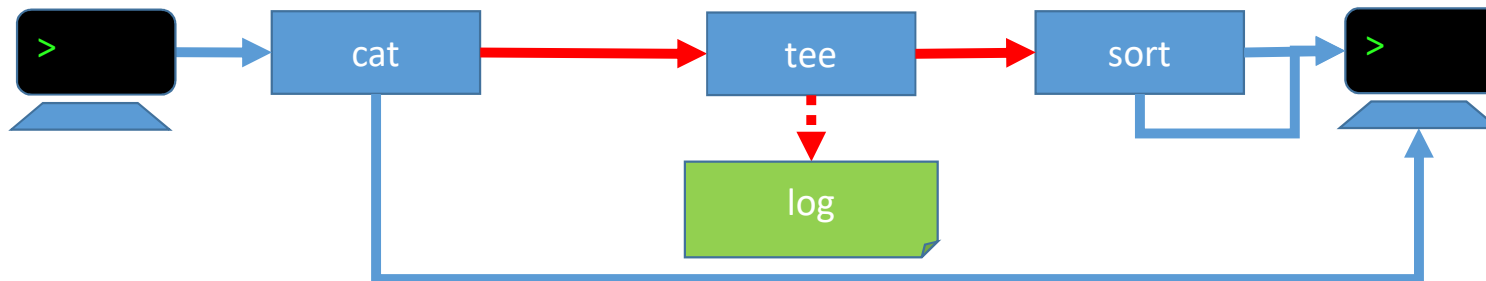
```
cat log
```

```
One  
Two  
Four
```

I don't always drink tea.



But when I do, I prefer
Earl Grey, hot.




How to pipe both STDOUT and STDERR?

- This is what happens with just a regular |

```
myprog.sh | tee output
This is something for STDOUT
This is something for STDERR

cat output
This is something for STDOUT
```




How to pipe both STDOUT and STDERR?

Not
working on
MacOS

```
myprog.sh |& tee output
This is something for STDOUT
This is something for STDERR

cat output
This is something for STDOUT
This is something for STDERR
```

- Now, we use the **|&** operator instead
- Both STDOUT and STDERR of myprog.sh were redirected to the STDIN of tee
- The STDOUT of tee, as well as the file, contain both messages



How do we **make sure** all ends up in STDOUT of tee?

```
myprog.sh |& tee output 2> /dev/null
This is something for STDOUT
This is something for STDERR

cat output
This is something for STDOUT
This is something for STDERR
```

- We redirect STDERR of tee to /dev/null
- We still get both msgs on screen
- Therefore, there was nothing on STDOUT coming out of tee

